

Shubhangi Jeve

B.Tech in Artificial Intelligence & Data Science
Terna College of Engineering, Dharashiv | CGPA: 8.5
Hyderabad, Telangana

+91-9579122372
shubhangijeve@gmail.com

shubhangijeve.github.io/My-Portfolio
github.com/ShubhangiJeve
linkedin.com/in/shubhangi-jeve-97445b235







CAREER OBJECTIVE

AI Engineer with 1.5+ years of core experience architecting and deploying production-grade AI systems, hybrid RAG pipelines, and multi-agent platforms. Proficient in LLM orchestration, async backend architectures, vector databases, and enterprise AI engineering. Driven to build scalable, high-impact intelligent systems.

EXPERIENCE

• AI Engineer Intern
COGNITBOTZ

November 2025 – Present
Hyderabad, Telangana

- Project 1: LegalAID** — Enterprise-grade AI Legal Research and Litigation Platform combining Hybrid RAG, semantic search, and automated drafting.
- **Hybrid RAG Architecture:** Engineered dense-lexical retrieval (pgvector BGE-M3 + Postgres FTS) with parent-child chunking, achieving 94% citation precision.
 - **LLM Gateway & Guardrails:** Architected multi-model LLM Gateway connected with OpenRouter and NVIDIA NIM with source validation and strict abstention thresholds, driving hallucination rates below 1%.
 - **High-Throughput Ingestion:** Automated ETL pipeline parsing 75,000+ court JSON cases over 25 years with schema normalization and deduplication in a single run.
 - **Async Backend API (FastAPI):** Developed async REST endpoints with Pydantic v2 validation for sub-4s semantic case search, RAG Q&A, and legal drafting.
 - **Full-Stack UI:** Built TypeScript workspace with natural language query search, faceted case explorer, drafting studio, and ingestion dashboard.
 - **Unified Vector DB (pgvector):** Leveraged PostgreSQL with pgvector for relational and ANN retrieval, eliminating external vector DB overhead and cutting latency by 45%.
 - **Production Containerization:** Containerized full stack via Docker Compose with health-check dependency chains, ensuring 99.9% uptime and data persistence.

- Project 2: MeetOps** — Enterprise AI Meeting Copilot and Project Intelligence Platform integrating Teams via ephemeral bots and agentic workflows.
- **Transcript RAG Pipeline:** Architected semantic RAG with speaker-turn chunking and pgvector embeddings, lifting topical retrieval precision by 35%.
 - **Confidence Abstention Gating:** Implemented 0.75 cosine similarity guardrail instructing LLMs to abstain on low context, eliminating fabricated technical details.
 - **Agentic Tool-Calling:** Engineered OpenAI function-calling workflow automating Jira ticket creation, rolling summaries, and live SAP GRC metrics retrieval.
 - **Disposable Bot Infrastructure:** Built ephemeral Dockerized Playwright bots with TTL watchdogs, guaranteeing zero state contamination across client meetings.
 - **Low-Latency Streaming:** Deployed multi-tier Redis semantic cache with SSE streaming, reducing perceived generation latency by 60% during live sessions.

- Project 3: Granite Buyer Intelligence** — Local-first B2B intelligence and lead generation platform automating buyer discovery, evidence qualification, and M365 outreach.
- **Modular Monolith Architecture:** Designed 9-domain lifecycle system (Discovery, Evidence, Intelligence, Email, Jobs) delivered solo from design to production.
 - **Multi-Tier LLM Failover Gateway:** Implemented 3-tier failover (OpenRouter Nemotron-550B → NVIDIA Build → Groq) and TF lexical search, guaranteeing 100% uptime.
 - **Durable Task Queue:** Built crash-resilient PostgreSQL job queue with 180s leases and heartbeat recovery, handling automated multi-engine web discovery.
 - **Evidence-First Validation:** Enforced SQL check constraints separating verified facts from AI inference, with 0–100 quality scoring and audit tracking.

• Artificial Intelligence Intern
Infosys Springboard

February 2025 – April 2025
Remote

- Built a Healthcare RAG Chatbot on a Wikipedia dataset using FAISS vector search, Flask backend, and React/Tailwind frontend for grounded clinical Q&A.
- Implemented a real-time query processing pipeline with source verification and an interactive UI supporting multi-turn question answering sessions.

• Data Science Intern
Adhyayan IT

January 2025 – June 2025
Remote

- Built a Text-to-SQL system using LangChain and LLMs (Gemini, Groq) with RAGAS evaluation achieving 100% context precision and high helpfulness scores.
- Developed an interactive Streamlit UI with secure MySQL connectivity, enabling non-technical users to query databases efficiently.

● **Data Analyst Intern**
Cyber Police Station

October 2023 – January 2024
Dharashiv, Maharashtra

- Managed 500+ user credentials for the National Cyber Crime Portal and organized 1,000+ complaints, improving case resolution efficiency by 30%.
- Communicated with victims to provide timely status updates, enhancing overall satisfaction by 15%.

PROJECTS

Project 1: Argus — Multi-Agent AI Platform

Autonomous hierarchical multi-agent platform with specialized tool routing, runtime guardrails, and full observability

- **Hierarchical Agent Routing:** Architected planner-specialist framework decomposing goals across research, analysis, and writing agents with scoped tool permissions.
- **Decoupled Async Backend (FastAPI):** Built async service routing requests through OpenRouter for dynamic task-based model selection and provider failover.
- **Dual-Ended Guardrail Pipeline:** Enforced input prompt-injection detection, strict tool allowlists, and Pydantic output schema validation with PII scrubbing.
- **Full Observability with Langfuse:** Instrumented end-to-end agent traces tracking latency, tool invocations, and token expenses to evaluate prompt revisions.
- **Real-Time Streaming UI & Azure Cloud:** Created React/TypeScript UI streaming agent step execution; deployed on Azure VM with Docker, Nginx, and TLS.

Project 2: LedgerLens-AI — Multi-LLM Document Intelligence & Forensic Engine

Production-grade multi-model orchestration engine leveraging Databricks Unity Catalog, FastAPI, and OpenRouter for high-throughput financial extraction and contract reconciliation

- **Multi-Model LLM Routing:** Engineered dynamic routing via OpenRouter directing forensic deduction to DeepSeek R1 and entity extraction to GPT-4o, cutting API inference costs by 42%.
- **Databricks Lakehouse & Unity Catalog:** Architected data pipelines on Databricks Unity Catalog and Delta Lake, enforcing data lineage, audit logging, and RBAC governance across 100,000+ financial documents.
- **Structured Output Validation & Caching:** Developed Pydantic schema-repair pipelines with PostgreSQL prompt-caching, achieving 99.4% adherence and reducing TTFT by 350ms.
- **Chain-of-Thought Forensics:** Implemented automated reasoning cross-examining balance sheets against vendor agreements, identifying \$1.8M+ in unapplied volume rebates.

Project 3: RAG Document QA System

Production-grade AI-powered question answering on PDF and TXT documents using FastAPI and React

- **Full-Stack Document QA:** Engineered FastAPI and React/TypeScript system with ChromaDB and sentence-transformers for verified Q&A across PDF/TXT files.
- **Dockerized API Deployment:** Built async REST APIs and deployed complete stack via Docker with an interactive dashboard for monitoring ingestion KPIs.

Project 4: Insurance Claim Automation System

AI-powered insurance claim processing using Google Gemini AI, LangChain, and Flask

- **Gemini-Powered Claim Extraction:** Developed automated medical claim validation pipeline using Google Gemini AI and NLP, reducing manual review time by 70%.
- **LangChain Decision Engine:** Built Flask web application with LangChain integration for intelligent claim assessment, error handling, and data validation.

TECHNICAL SKILLS

AI / LLMs: Large Language Models (OpenAI, DeepSeek R1, Claude, Azure OpenAI, Llama, Gemini), Fine-Tuning & Integration, LangChain, LangGraph, AI Orchestration, Multi-Agent Architectures, Prompt Engineering, Prompt-Caching & Registry

RAG & Search: Retrieval-Augmented Generation (RAG), Semantic Search, Azure AI Search, Weaviate, FAISS, pgvector (Azure PostgreSQL), ChromaDB

Backend & APIs: Python, FastAPI, Flask, REST APIs & Microservices for AI Applications, Pydantic, Structured Outputs, PostgreSQL, PgSQL, Async Programming

MLOps & Cloud: Databricks (Unity Catalog, Delta Lake, PySpark, MLflow), MLOps (Model Deployment, Monitoring, Evaluation & Quality Tuning), Containerization (Docker, Kubernetes), Azure Cloud Platform, CI/CD, Git, GitLab

EDUCATION

● **Bachelor of Technology in Artificial Intelligence and Data Science**
Terna Public Charitable Trust's College of Engineering, Dharashiv

2021 – 2025
CGPA: 8.5